

Performance Optimization Report

Introduction

The task involved profiling, optimizing, and analyzing two matrix-vector multiplication functions: `my_dense` for dense matrices for sparse matrices. The objective was to improve execution time by leveraging compiler optimizations, vectorization, parallelism, and memory access patterns. The initial implementations were benchmarked, and subsequent optimizations were applied to enhance performance. This report documents the analysis, optimizations applied, and results obtained.

Original Execution Time

The performance of the original implementations for the matrix-vector multiplication functions were measured using execution time metrics:

- **Original `my_dense` Execution Time:** 710 ms (GCC O0)

Optimized Execution Time

Optimizations were applied to both dense and sparse matrix-vector multiplication functions. The optimized execution times for the `my_dense` kernel are presented in the following table for different compiler optimization flags:

`my_dense` Execution Time (in ms):

Compiler	O0	O2	O3-VEC	OFAST-VEC	REF
GSL(GCC)	710	460	230.66	210.66	328.67
MKL(ICC)	730.66	480.66	270.65	230.33	345.66

Note: The optimized `my_dense` kernel shows improvements across various compiler optimizations, with the best performance achieved using `Ofast-vec`.

Tools and Methods Used

To profile, optimize, and analyze the performance of the matrix-vector multiplication implementations, the following tools were employed:

- **gprof (GNU Profiler):** Used for profiling the original and optimized implementations, identifying performance bottlenecks and hotspots in both functions.
- **OpenMP:** Applied to parallelize the matrix-vector multiplication operations. OpenMP directives were used to parallelize outer loops, enhancing performance

by utilizing multiple CPU cores.

- **Intel MKL:** Used for optimized linear algebra routines that contributed to improved performance, particularly in the dense matrix multiplication (`my_dense`).
- **Valgrind:** Utilized to ensure efficient memory management, detect memory leaks, and verify correct memory access patterns during optimization.
- **Perf:** Used for low-level performance analysis, including CPU cycles, cache behavior, and instruction counts. Perf highlighted cache misses and CPU cycles, guiding further optimizations.

Analysis of Vectorization and Memory/Cache Behavior

Vectorization and Parallelization

- **GCC Optimizations:** The `O3-vec` and `Ofast-vec` flags enabled compiler vectorization. These optimizations helped improve the performance of the `my_dense` kernel by enabling the use of vector instructions, allowing for multiple data elements to be processed simultaneously.
- **Intel MKL Optimizations:** The use of Intel MKL routines enhanced performance by leveraging highly optimized libraries for matrix multiplication, which take advantage of hardware-level vectorization and parallelization.

Memory and Cache Optimization

- **Memory Efficiency:** Memory allocation was optimized to reduce overhead and avoid memory leaks. The use of Valgrind ensured that the optimizations did not introduce memory issues.
- **Cache Utilization:** Manual loop unrolling and optimizations to improve spatial and temporal locality led to better cache utilization. Cache-friendly access patterns, particularly in the sparse matrix implementation, reduced cache misses and enhanced performance.

Speedup Results and Discussion

Dense Matrix Optimization (`my_dense`)

- **GSL (GCC) Optimized Execution Time:**
 - Original: 710 ms (O0)
 - Optimized (`Ofast-vec`): 210.66 ms
 - **Speedup:** The optimized version achieved a **70.3%** reduction in execution time.
- **Intel MKL (ICX) Optimized Execution Time:**
 - Original: 730.66 ms (O0)
 - Optimized (`Ofast-vec`): 230.33 ms
 - **Speedup:** The optimized version achieved a **68.4%** reduction in execution

time.

Conclusion

The optimizations applied to both the dense and sparse matrix-vector multiplication functions resulted in significant performance improvements:

- **Dense Matrix (my_dense):** The optimized version showed a **70.3%** performance improvement using GCC and a **68.4%** improvement using Intel MKL. These gains were achieved through compiler optimizations, vectorization, and parallelization.

Profiling tools like gprof, Valgrind, and Perf were instrumental in identifying bottlenecks and verifying the correctness of the optimizations. The improvements demonstrate that with careful optimization of memory access patterns, vectorization, and parallelism, substantial performance gains can be achieved, especially for larger matrices. The results indicate that the optimized implementations are suitable for real-world applications involving large-scale matrix operations.